



Assembly 101

x86/x64

EPITA · APPING2 S7

Auteur(s) Hugo Sibony

Date 24/12/2025

Table des matières

| | | |
|-------|--|----|
| 1 | What is Assembly? | 3 |
| 1.1 | Simple Example | 3 |
| 1.2 | Structure of an Assembly File | 3 |
| 2 | Registers: Your Fast Memory | 4 |
| 2.1 | 64-bit Registers (amd64) | 4 |
| 2.2 | 32-bit Registers (i386) | 4 |
| 2.3 | Register Sub-parts | 5 |
| 3 | The Two Syntaxes: AT&T vs Intel | 6 |
| 3.1 | Comparison Table | 6 |
| 3.2 | Side-by-Side Examples | 6 |
| 3.3 | Size Suffixes (AT&T) | 7 |
| 3.4 | Intel Template with GAS | 7 |
| 4 | Loops and Conditions | 8 |
| 4.1 | CPU Flags | 8 |
| 4.2 | Setting a Register to Zero | 8 |
| 4.3 | Comparing Two Values: <code>cmp</code> | 8 |
| 4.4 | Testing if a Register is Zero: <code>test</code> | 8 |
| 4.5 | Conditional Jumps | 9 |
| 4.6 | Example: Simple Condition | 9 |
| 4.7 | Example: While Loop | 10 |
| 4.8 | Example: Traversing a String | 10 |
| 4.9 | Example: For Loop with Array | 11 |
| 4.10 | Tip: <code>lea</code> for Arithmetic | 12 |
| 5 | Division | 13 |
| 5.1 | 64-bit Division | 13 |
| 5.2 | Signed Division | 13 |
| 6 | The Stack: How It Works | 14 |
| 6.1 | Basic Operations | 14 |
| 6.2 | Stack Visualization After a CALL | 14 |
| 7 | Stack Frame and Alignment | 15 |
| 7.1 | Why 16 Bytes? | 15 |
| 7.1.1 | What is SSE? | 15 |
| 7.1.2 | Why Alignment? | 15 |
| 7.2 | The Mechanism | 15 |
| 7.3 | How to Realign | 15 |
| 7.3.1 | Solution 1: Single Push | 16 |
| 7.3.2 | Solution 2: Direct Sub | 16 |
| 7.3.3 | Solution 3: Frame Pointer | 16 |
| 7.4 | Visual Summary | 16 |
| 7.5 | Complete Stack Frame | 17 |
| 7.6 | Common Mistake | 17 |
| 8 | The ABI: The Contract Between Functions | 18 |
| 8.1 | i386 ABI (32-bit) | 18 |

| | | |
|------|---|----|
| 8.2 | amd64 ABI (64-bit) | 18 |
| 8.3 | Saving Registers: When and How? | 18 |
| 9 | Calling C Functions from Assembly | 19 |
| 9.1 | Recipe | 19 |
| 9.2 | Stack Alignment | 19 |
| 9.3 | @PLT and @GOTPCREL | 19 |
| 9.4 | lea vs mov | 20 |
| 9.5 | Example: Calling puts | 20 |
| 9.6 | printf with Multiple Arguments | 20 |
| 9.7 | Libc File Functions | 20 |
| 10 | Syscalls: Talking to the Linux Kernel | 21 |
| 10.1 | Difference from C Functions | 21 |
| 10.2 | Syscall Structure | 21 |
| 10.3 | Numbers and Arguments | 21 |
| 10.4 | Standard File Descriptors | 22 |
| 10.5 | Flags for open() | 22 |
| 10.6 | Flags for mmap() | 22 |
| 10.7 | mmap() to Allocate Memory | 23 |
| 10.8 | Error Handling | 23 |
| 11 | Debugging with GDB | 24 |
| 11.1 | Launching GDB | 24 |
| 11.2 | Example Session | 24 |
| 12 | Quick Reference | 25 |
| 12.1 | Calling Conventions | 25 |
| 12.2 | Callee-saved (must save) | 25 |
| 12.3 | Syscalls | 25 |
| 12.4 | Templates | 25 |

1 What is Assembly?

Assembly is the language closest to the processor. Each instruction corresponds directly to a CPU operation.

1.1 Simple Example

In C:

```
int x = 5;
x = x + 3;
```

In assembly, we directly manipulate CPU registers:

```
movl $5, %eax      # Put 5 in register EAX
addl $3, %eax      # Add 3 to EAX → EAX now contains 8
```

1.2 Structure of an Assembly File

```
.section .data          # Data section (global variables)
message: .asciz "Hello" # A string

.section .text           # Code section
.globl my_function        # Make the symbol visible to the linker

my_function:             # Label = start of function
    # your code here
    ret                   # Return to caller
```

* Important Directives

| Directive | Purpose |
|------------------|---|
| .section .text | Executable code |
| .section .data | Initialized data (modifiable) |
| .section .rodata | Read-only data (constants, strings) |
| .globl name | Export the symbol (required for C to find your function!) |
| .ascii "..." | String without trailing \0 |
| .asciz "..." | String with trailing \0 |
| .long value | 32-bit value |
| .quad value | 64-bit value |

2 Registers: Your Fast Memory

Registers are « variables » directly inside the CPU. They are ultra-fast.

2.1 64-bit Registers (amd64)

| Register | Purpose |
|----------|---|
| RAX | Function return value |
| RBX | General purpose (MUST SAVE!) |
| RCX | 4th argument (historically: loop counter in i386) |
| RDX | 3rd argument / used by divisions |
| RSI | 2nd argument |
| RDI | 1st argument |
| RSP | Stack pointer (NEVER LOSE IT!) |
| RBP | Base pointer (MUST SAVE!) |
| R8 | 5th argument |
| R9 | 6th argument |
| R10-R11 | Temporaries |
| R12-R15 | General purpose (MUST SAVE!) |

2.2 32-bit Registers (i386)

| Register | Purpose |
|----------|------------------------------|
| EAX | Return value |
| EBX | General purpose (MUST SAVE!) |
| ECX | Counter (for loops) |
| EDX | Data / used by divisions |
| ESI | Source (MUST SAVE!) |
| EDI | Destination (MUST SAVE!) |
| ESP | Stack pointer |
| EBP | Base pointer (MUST SAVE!) |

2.3 Register Sub-parts

A 64-bit register can be accessed in parts:



! Important

WARNING: Writing to EAX automatically zeroes the upper 32 bits of RAX!

```
movq $0xFFFFFFFFFFFFFFF, %rax    # RAX = 0xFFFFFFFFFFFFFFF  
movl $1, %eax                   # RAX = 0x0000000000000001 (not 0xFFFFFFFF00000001!)
```

3 The Two Syntaxes: AT&T vs Intel

You need to know how to write in both syntaxes. Here are the differences:

3.1 Comparison Table

| Aspect | AT&T | Intel |
|------------|---|--|
| Order | source, destination | destination, source |
| Registers | %rax | rax |
| Immediates | \$42 | 42 |
| Memory | 8(%rsp) | [rsp + 8] |
| Size | Suffix: <code>movl</code> , <code>movq</code> | Keyword: <code>DWORD</code> , <code>QWORD</code> |

3.2 Side-by-Side Examples

Put a value in a register:

AT&T

```
movl $42, %eax
movq $100, %rax
```

Intel

```
mov eax, 42
mov rax, 100
```

Copy one register to another:

AT&T

```
movq %rdi, %rax
```

Intel

```
mov rax, rdi
```

Read from memory (stack):

AT&T

```
movl 4(%esp), %eax
```

Intel

```
mov eax, [esp + 4]
```

Write to memory:

AT&T

```
movq %rax, (%rdi)
movq %rax, 8(%rdi)
```

Intel

```
mov [rdi], rax
mov [rdi + 8], rax
```

3.3 Size Suffixes (AT&T)

| Suffix | Size | Example Register |
|--------|----------------------|-----------------------|
| b | 1 byte | AL , BL , CL , R8B |
| w | 2 bytes (word) | AX , BX , CX , R8W |
| l | 4 bytes (long/dword) | EAX , EBX , ECX , R8D |
| q | 8 bytes (quad) | RAX , RBX , RCX , R8 |

3.4 Intel Template with GAS

To use Intel syntax with the GNU Assembler (GAS):

```
.intel_syntax noprefix      # Enable Intel, disable % and $ prefixes

.text
.globl my_function
my_function:
    # your code in Intel syntax here
    ret
```

4 Loops and Conditions

4.1 CPU Flags

The CPU has « flags » updated by certain instructions:

- **ZF** (Zero Flag): set to 1 if the result is zero
- **SF** (Sign Flag): set to 1 if the result is negative

Jump instructions (`jz`, `jnz`, etc.) read these flags to decide whether to jump.

4.2 Setting a Register to Zero

```
xorq %rax, %rax           # RAX = RAX XOR RAX = 0
```

* Astuce

Why `xor` and not `movq $0, %rax` ?

- **Shorter:** `xorl %eax, %eax` is 2 bytes vs 5 bytes for `movl $0, %eax`
- **Faster:** CPUs recognize this pattern and optimize it (zero-idiom)
- **Bonus:** Writing to EAX automatically zeros the upper 32 bits of RAX!

4.3 Comparing Two Values: `cmp`

```
cmpq %rsi, %rdi           # Computes RDI - RSI, sets flags, discards result
```

After `cmpq %rsi, %rdi` :

- If `RDI == RSI` → `ZF = 1`
- If `RDI < RSI` → `SF = 1` (for signed numbers)

4.4 Testing if a Register is Zero: `test`

```
testq %rdi, %rdi           # Computes RDI AND RDI, sets flags, discards result
```

`RDI AND RDI` = RDI. Therefore:

- If `RDI == 0` → `result = 0` → `ZF = 1`
- If `RDI != 0` → `result != 0` → `ZF = 0`

This is the standard way to test if a register is zero.

4.5 Conditional Jumps

| Instruction | Condition (after <code>cmp b, a</code> or <code>test</code>) |
|------------------------|---|
| <code>je / jz</code> | <code>a == b</code> (or result == 0) |
| <code>jne / jnz</code> | <code>a != b</code> (or result != 0) |
| <code>jl / jg</code> | <code>a < b</code> / <code>a > b</code> (signed) |
| <code>jle / jge</code> | <code>a <= b</code> / <code>a >= b</code> (signed) |
| <code>jb / ja</code> | <code>a < b</code> / <code>a > b</code> (unsigned) |
| <code>jbe / jae</code> | <code>a <= b</code> / <code>a >= b</code> (unsigned) |
| <code>js</code> | result is negative |
| <code>jmp</code> | always |

★ Astuce

Signed vs unsigned: `jl / jg` for int, `jb / ja` for unsigned

4.6 Example: Simple Condition

In C:

```
if (x == 0) return -1;
return x;
```

In assembly:

```
testq %rdi, %rdi          # RDI AND RDI → ZF=1 if RDI==0
jnz not_zero               # jump if ZF=0 (so if RDI ≠ 0)
movq $-1, %rax
ret
not_zero:
    movq %rdi, %rax
    ret
```

4.7 Example: While Loop

In C:

```
int sum = 0;
while (n > 0) { sum += n; n--; }
return sum;
```

In assembly:

```
xorq %rax, %rax          # sum = 0 (xor with itself = 0)
loop:
    testq %rdi, %rdi      # n == 0?
    jz done                # if ZF=1 (n==0), finish
    addq %rdi, %rax        # sum += n
    decq %rdi              # n--
    jmp loop
done:
    ret
```

4.8 Example: Traversing a String

In C:

```
int len = 0;
while (*s != '\0') { len++; s++; }
return len;
```

In assembly:

```
xorq %rax, %rax          # len = 0
loop:
    movzbl (%rdi), %ecx    # load 1 byte at address RDI into ECX
    testb %cl, %cl         # byte == 0?
    jz done                # if ZF=1 (byte==0), finish
    incq %rax              # len++
    incq %rdi              # s++ (move to next character)
    jmp loop
done:
    ret
```

i Information

`movzbl (%rdi), %ecx` = «Move Zero-extend Byte to Long»

- Reads 1 byte at address RDI
- Zero-extends it to 32 bits
- Puts it in ECX

Intel equivalent: `movzx ecx, BYTE PTR [rdi]`

4.9 Example: For Loop with Array

In C:

```
int sum = 0;
for (int i = 0; i < len; i++) { sum += arr[i]; }
return sum;
```

In assembly:

```
xorl %eax, %eax      # sum = 0
xorl %ecx, %ecx      # i = 0
loop:
    cmpl %esi, %ecx      # compare i with len
    jge done              # if i ≥ len, finish
    movl (%rdi,%rcx,4), %edx  # edx = arr[i]
    addl %edx, %eax      # sum += arr[i]
    incl %ecx            # i++
    jmp loop
done:
    ret
```

* Indexed Addressing

(%rdi,%rcx,4) = address RDI + RCX * 4

- RDI = base address of array
- RCX = index
- 4 = element size (4 bytes for `int`)

4.10 Tip: lea for Arithmetic

`lea` (Load Effective Address) computes an address without accessing memory. It's a powerful tool for arithmetic!

* Why use `lea` instead of `add/imul`?

- **Doesn't modify flags:** You can keep flags from a previous `cmp / test`
- **Multiple operations in one:** `lea (%rdi,%rsi,4), %rax` does `rax = rdi + rsi*4` in one instruction
- **Result in different register:** `lea 1(%rdi), %rax` computes `rdi + 1` into `rax` without modifying `rdi`

AT&T

```
leaq (%rdi,%rsi), %rax    # rax = rdi + rsi
leaq (%rdi,%rdi,4), %rax  # rax = rdi * 5
leaq 1(%rdi), %rax        # rax = rdi + 1
```

Intel

```
lea rax, [rdi + rsi]
lea rax, [rdi + rdi*4]
lea rax, [rdi + 1]
```

* Multiplication by 10 (useful for atoi)

To compute `result = result * 10 + digit` (where digit is in RDX):

```
leaq (%rax,%rax,4), %rax      # rax = rax + rax*4 = rax * 5
leaq (%rdx,%rax,2), %rax      # rax = rdx + rax*2 = digit + (old_rax*5)*2
                                #
                                = digit + old_rax*10
```

Step by step: If RAX=42 and RDX=7:

- After 1st `lea`: $RAX = 42 + 42 \times 4 = 210$ (which is 42×5)
- After 2nd `lea`: $RAX = 7 + 210 \times 2 = 427$ (which is $7 + 42 \times 10$)

5 Division

Division in assembly is a bit special because it uses a pair of registers.

5.1 64-bit Division

```
xorq %rdx, %rdx      # REQUIRED! Set RDX to 0
divq %rbx             # Divide RDX:RAX by RBX
                      # RAX = quotient
                      # RDX = remainder
```

! Important

`xorq %rdx, %rdx` is **mandatory** before `divq`. The instruction divides the 128-bit number formed by RDX:RAX. If RDX contains random values, the result will be wrong.

5.2 Signed Division

For signed numbers, use `idivq` and `cqto` (sign-extend RAX into RDX:RAX):

```
cqto
idivq %rbx          # Sign-extend RAX into RDX
                      # Signed division
```

6 The Stack: How It Works

The stack is a memory area that **grows downward** (decreasing addresses).

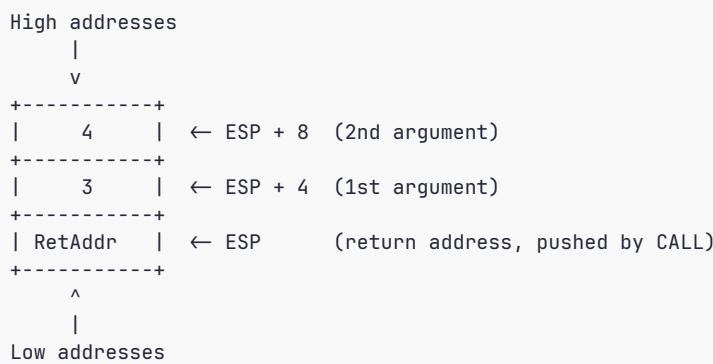
6.1 Basic Operations

```
# PUSH: Push a value onto the stack
pushq %rax
# Equivalent to:
#   1. RSP = RSP - 8      (stack grows downward)
#   2. Write RAX to address RSP

# POP: Pop a value from the stack
popq %rax
# Equivalent to:
#   1. Read value at address RSP into RAX
#   2. RSP = RSP + 8      (stack shrinks)
```

6.2 Stack Visualization After a CALL

When C calls your function `add(3, 4)` in 32-bit:



! Important

This is why the 1st argument is at ESP+4 and not ESP+0! The `CALL` instruction automatically pushes the return address.

7 Stack Frame and Alignment

7.1 Why 16 Bytes?

The ABI requires: **RSP % 16 == 0 before each `call`**.

7.1.1 What is SSE?

SSE (Streaming SIMD Extensions) = instructions that process **multiple data in parallel**.

The CPU has special 128-bit (16-byte) registers: `XMM0` to `XMM15`.

```
# Classic registers: 1 operation at a time
addq %rax, %rbx      # 1 64-bit addition

# SSE registers: 4 operations in parallel
addps %xmm0, %xmm1    # 4 32-bit additions at once
```

The libc uses SSE to optimize `memcpy`, `strlen`, `printf`, etc. It's faster.

7.1.2 Why Alignment?

SSE instructions that access memory (`movaps`, `movdqa`) require an address divisible by 16. This is a hardware constraint: the memory bus loads 16 aligned bytes more efficiently.

```
movaps (%rsp), %xmm0      # Load 16 bytes from RSP into XMM0
                            # If RSP % 16 ≠ 0 → immediate SIGSEGV
```

If your stack is misaligned when you call `printf` → `printf` uses `movaps` → crash.

7.2 The Mechanism

Let's follow RSP step by step with real values:

1. Someone calls your function
RSP = 0xffff0010 (multiple of 16 ✓)
2. The CALL instruction executes
 - Pushes return address (8 bytes)
 - RSP = 0xffff0010 - 8 = 0xffff0008
3. You enter your function
RSP = 0xffff0008 → 0xffff0008 % 16 = 8 × MISALIGNED

Problem: At the entry of your function, RSP is **ALWAYS** misaligned by 8.

7.3 How to Realign

You must subtract 8 bytes (or an odd number of times 8) to get back to a multiple of 16.

7.3.1 Solution 1: Single Push

```
func:
    pushq %rbx          # RSP -= 8
    # RSP was 0xfffff0008, now 0xfffff0000 ✓

    call printf@PLT      # OK, RSP aligned

    popq %rbx
    ret
```

Calculation: $0xfffff0008 - 8 = 0xfffff0000 \rightarrow 0xfffff0000 \% 16 = 0 \checkmark$

7.3.2 Solution 2: Direct Sub

```
func:
    subq $8, %rsp        # Same effect as a push

    call printf@PLT

    addq $8, %rsp
    ret
```

7.3.3 Solution 3: Frame Pointer

```
func:
    pushq %rbp          # RSP -= 8 → aligned
    movq %rsp, %rbp
    subq $32, %rsp       # 32 bytes of local variables (multiple of 16)

    # Variables: -8(%rbp), -16(%rbp), -24(%rbp), -32(%rbp)

    leave                # RSP = RBP, pop RBP
    ret
```

Why `sub $32` and not `sub $24` if I have 3 variables of 8 bytes?

```
After push %rbp: RSP = 0xfffff0000 (aligned)
After sub $24:   RSP = 0xfffff0000 - 24 = 0x7ffefff8
                 0x7ffefff8 \% 16 = 8 ×

After sub $32:   RSP = 0xfffff0000 - 32 = 0x7ffe0000
                 0x7ffe0000 \% 16 = 0 ✓
```

Rule: after `push %rbp`, always allocate a multiple of 16.

7.4 Visual Summary

| RSP \% 16 | | |
|--------------------|---|--------------|
| <hr/> | | |
| Before call | 0 | ✓ aligned |
| After call (entry) | 8 | ✗ misaligned |
| After 1 push | 0 | ✓ aligned |
| After 2 push | 8 | ✗ misaligned |
| After 3 push | 0 | ✓ aligned |

Pattern: odd number of pushes → aligned. Even number → misaligned.

7.5 Complete Stack Frame

```
+-----+
+24 | argument 7      | (if > 6 args)
+-----+
+16 | argument 8      |
+-----+
+8  | return address   | ← pushed by CALL
+-----+
0   | old RBP          | ← RBP points here
+-----+
-8  | variable 1       |
+-----+
-16 | variable 2       | ← RSP after sub $16
+-----+
```

Access: `movq -8(%rbp), %rax` to read variable 1.

7.6 Common Mistake

```
# WRONG - will segfault
func:
    call printf@PLT      # RSP = 8 mod 16 → CRASH
    ret

# CORRECT
func:
    pushq %rbx
    call printf@PLT      # RSP = 0 mod 16 → OK
    popq %rbx
    ret
```

! Important

Debug: In GDB, before a suspicious call: `p $rsp % 16`

If it shows 8 → your stack is misaligned.

8 The ABI: The Contract Between Functions

The ABI (Application Binary Interface) defines the rules for functions to call each other correctly.

8.1 i386 ABI (32-bit)

| Category | Value |
|--------------|--|
| Arguments | On stack: [ESP+4] , [ESP+8] , [ESP+12] ... |
| Return | EAX |
| Caller-saved | EAX, ECX, EDX |
| Callee-saved | EBX, ESI, EDI, EBP |

8.2 amd64 ABI (64-bit)

| Category | Value |
|--------------|---|
| Arguments | RDI, RSI, RDX, RCX, R8, R9 (in this order) |
| Return | RAX |
| Caller-saved | RAX, RCX, RDX, RSI, RDI, R8-R11 |
| Callee-saved | RBX, RBP, R12-R15 |
| Stack | RSP multiple of 16 before <code>call</code> |

8.3 Saving Registers: When and How?

If your function uses a « callee-saved » register (like RBX), you MUST:

1. Save it at the beginning (with push)
2. Restore it at the end (with pop)

```
my_function:
    pushq %rbx           # Save RBX
    # ... you can use RBX here ...
    popq %rbx            # Restore RBX
    ret
```

If you only use « caller-saved » registers (RAX, RDI, RSI...), no need to save them.

9 Calling C Functions from Assembly

9.1 Recipe

1. Put arguments in RDI, RSI, RDX, RCX, R8, R9
2. Align the stack (RSP must be multiple of 16 before `call`)
3. `call function@PLT` → result in RAX

9.2 Stack Alignment

At entry of your function, $RSP = 8 \bmod 16$ (return address was pushed).

Simple solution: `pushq %rbx` before `call` (adds 8 bytes → RSP multiple of 16).

Solution with frame pointer (recommended for complex functions):

AT&T:

```
my_function:
    pushq %rbp
    movq %rsp, %rbp
    subq $16, %rsp

    # ... your code ...
    # Local variables: -8(%rbp), -16(%rbp)

    leave
    ret
```

Intel:

```
my_function:
    push rbp
    mov rbp, rsp
    sub rsp, 16

    ; ... your code ...
    ; Local variables: [rbp-8], [rbp-16]

    leave
    ret
```

* Astuce

`leave` = restore RSP from RBP then pop RBP. Handy for cleaning up the stack frame.

9.3 @PLT and @GOTPCREL

When calling libc functions or accessing external global variables:

```
# Call an external function (printf, puts, getline, etc.)
call printf@PLT          # @PLT = Procedure Linkage Table

# Access an external global variable (stdin, stdout, stderr)
movq stdin@GOTPCREL(%rip), %rdx  # Load ADDRESS of stdin from GOT
movq (%rdx), %rdx               # Load VALUE of stdin (the FILE*)
```

i Information

Why @PLT? The dynamic linker resolves the function address at runtime.

Why @GOTPCREL? External global variables are in the Global Offset Table.

9.4 lea vs mov

```
leaq msg(%rip), %rdi      # RDI = ADDRESS of msg (to pass a string)
movq var(%rip), %rdi      # RDI = CONTENTS of var (to read a variable)
```

9.5 Example: Calling puts

```
.section .rodata
msg: .asciz "Hello"

.section .text
.globl example
example:
    pushq %rbx                # Align stack
    leaq msg(%rip), %rdi        # arg1 = address of msg
    call puts@PLT
    popq %rbx
    ret
```

9.6 printf with Multiple Arguments

```
leaq format(%rip), %rdi      # arg1: format string
movq %rcx, %rsi              # arg2: value to display
xorl %eax, %eax              # RAX = 0 (no floating-point arguments)
call printf@PLT
```

! Important

`xorl %eax, %eax` is mandatory before printf/scanf (variadic functions).

9.7 Libc File Functions

Functions for file manipulation:

| Function | Prototype | Description |
|----------|-------------------------------------|---|
| fopen | FILE *fopen(char *path, char *mode) | Open a file, returns NULL on error |
| fclose | int fclose(FILE *stream) | Close a file |
| fgetc | int fgetc(FILE *stream) | Read a character, returns -1 (EOF) at end |
| fputc | int fputc(int c, FILE *stream) | Write a character |
| putchar | int putchar(int c) | Write a character to stdout |

* Astuce

fopen modes: "r" = read, "w" = write (creates/truncates)

10 Syscalls: Talking to the Linux Kernel

10.1 Difference from C Functions

! Important

WARNING: For syscalls, the 4th argument is in **R10**, not RCX! (The `syscall` instruction uses RCX internally to save RIP)

| C Functions | Syscalls |
|-------------|------------------------|
| 4th arg | R10 |
| Call | <code>call func</code> |
| Error | Negative RAX |

10.2 Syscall Structure

```

movq $NUMBER, %rax      # Syscall number
movq arg1, %rdi        # 1st argument
movq arg2, %rsi        # 2nd argument
movq arg3, %rdx        # 3rd argument
movq arg4, %r10         # 4th argument (WARNING: R10, not RCX!)
movq arg5, %r8          # 5th argument
movq arg6, %r9          # 6th argument
syscall                # Execute syscall
# Result in RAX (negative = error)

```

10.3 Numbers and Arguments

| # | Name | Arguments |
|---|-------|--|
| 0 | read | fd, buffer, count → bytes read |
| 1 | write | fd, buffer, count → bytes written |
| 2 | open | path, flags, mode → fd |
| 3 | close | fd |
| 9 | mmap | addr, len, prot, flags, fd, offset → ptr |

10.4 Standard File Descriptors

| Name | Value | Description |
|--------|-------|-----------------|
| stdin | 0 | standard input |
| stdout | 1 | standard output |
| stderr | 2 | standard error |

10.5 Flags for open()

| Constant | Value | Description |
|----------|-------|-------------------------|
| O_RDONLY | 0 | Read only |
| O_WRONLY | 1 | Write only |
| O_RDWR | 2 | Read and write |
| O_CREAT | 0x40 | Create if doesn't exist |
| O_TRUNC | 0x200 | Truncate if exists |

★ Astuce

To open for writing with creation: O_WRONLY | O_CREAT | O_TRUNC = 0x241

10.6 Flags for mmap()

| Constant | Value |
|---------------|-------|
| PROT_READ | 1 |
| PROT_WRITE | 2 |
| MAP_PRIVATE | 0x02 |
| MAP_ANONYMOUS | 0x20 |

★ Astuce

To allocate memory:

- prot = 3 (PROT_READ | PROT_WRITE = 1 | 2)
- flags = 0x22 (MAP_PRIVATE | MAP_ANONYMOUS = 0x02 | 0x20)
- fd = -1 (no file, anonymous mapping)
- offset = 0

10.7 mmap() to Allocate Memory

```
movq $9, %rax          # syscall mmap
xorq %rdi, %rdi        # addr = NULL
movq $SIZE, %rsi        # len = desired size
movq $3, %rdx           # prot = PROT_READ | PROT_WRITE
movq $0x22, %r10         # flags = MAP_PRIVATE | MAP_ANONYMOUS
movq $-1, %r8            # fd = -1
xorq %r9, %r9           # offset = 0
syscall
# RAX = pointer to allocated memory (or negative if error)
```

10.8 Error Handling

Syscalls return a negative value on error:

```
syscall
testq %rax, %rax      # SF=1 if RAX negative
js error               # Jump if Sign (if negative)
```

11 Debugging with GDB

11.1 Launching GDB

```
gdb ./my_program
```

| Command | Description |
|----------------|----------------------------------|
| break add | Breakpoint on function add |
| run | Start |
| stepi | 1 instruction (steps into calls) |
| nexti | 1 instruction (steps over calls) |
| continue | Continue |
| info registers | View registers |
| p/x \$rax | Print RAX in hex |
| x/s \$rdi | Print string pointed to by RDI |
| x/10x \$rsp | 10 hex values from RSP |

11.2 Example Session

```
$ gdb ./my_program
(gdb) break add           ← breakpoint on your function
(gdb) run                 ← start the program
Breakpoint 1, add ()
(gdb) info registers rdi rsi ← check arguments
rdi          0x3
rsi          0x4
(gdb) stepi               ← execute 1 instruction
(gdb) info registers rax   ← check result
(gdb) continue            ← continue execution
```

12 Quick Reference

12.1 Calling Conventions

| Type | Arguments | Return |
|---------------------|--|--------|
| C functions (amd64) | RDI, RSI, RDX, RCX, R8, R9 | RAX |
| Syscalls (amd64) | RDI, RSI, RDX, R10 , R8, R9 (number in RAX) | RAX |
| i386 | [ESP+4], [ESP+8], ... | EAX |

12.2 Callee-saved (must save)

| Arch | Registers |
|-------|--------------------|
| amd64 | RBX, RBP, R12-R15 |
| i386 | EBX, ESI, EDI, EBP |

12.3 Syscalls

| # | Name | Args |
|---|-------|---------------------------------|
| 0 | read | fd, buf, count |
| 1 | write | fd, buf, count |
| 2 | open | path, flags, mode |
| 3 | close | fd |
| 9 | mmap | addr, len, prot, flags, fd, off |

12.4 Templates

amd64 AT&T

amd64 Intel

i386

| | | |
|--|---|--|
| <pre>.text .globl func func: ret</pre> | <pre>.intel_syntax noprefix .text .globl func func: ret</pre> | <pre>.text .globl func func: ret</pre> |
|--|---|--|